

Last Time: Direct / Pattern Search

- Coordinate Search
- Generalized Pattern Search
- Nelder-Mead Simplex Method
- DIRECT (Lipschitz-Inspired)



Today: Stochastic + Population Methods

- Simulated Annealing
- Cross-Entropy

2 challenges

1. Finding local min within neighborhood

2. Finding neighborhood with lowest min

so far

Simulated Annealing

temperature decay schedule

initialize x

hyperparameters $D, t(k)$

loop


$$x' \leftarrow x + \text{rand}(D)$$

$$\Delta y \leftarrow f(x') - f(x)$$

$$\text{accept} \leftarrow \text{rand}(U(0,1)) \leq \begin{cases} 1 & \text{if } \Delta y \leq 0 \\ e^{-\Delta y/t(k)} & \text{if } \Delta y > 0 \end{cases}$$

if accept

$$x \leftarrow x'$$



Common temperature schedules

$$t^k = t' \ln(z) / \ln(k+1) \quad \text{"logarithmic"}$$

asymptotically optimal, very slow

$$t^{k+1} = \gamma t^k \quad \text{"exponential"}$$

$$t^k = \frac{t'}{k} \quad \text{"fast annealing"}$$

Cross-Entropy Method

Parameterized Distribution $D(\theta)$

initialize $\theta \leftarrow$ Gaussian: μ, Σ

hyperparameters: D, m, m_{elite}

loop

population \leftarrow sample m from $D(\theta)$


elite \leftarrow select m_{elite} best samples

$\theta \leftarrow$ fit to elite R_{minimize} "cross entropy"

fit: If D is Gaussian $\Theta = (\mu, \Sigma)$

$$\vec{\mu}^{k+1} \leftarrow \frac{1}{m_{\text{elite}}} \sum_{i=1}^{m_{\text{elite}}} \vec{x}^i$$

$$\Sigma^{k+1} \leftarrow \frac{1}{m_{\text{elite}}} \sum_{i=1}^{m_{\text{elite}}} (\vec{x}^i - \vec{\mu}^{k+1})(\vec{x}^i - \vec{\mu}^{k+1})^T$$



Other Stochastic Algorithms



- Mesh-adaptive direct search \leftarrow stochastic pattern search

- Natural Evolution Strategies

estimate $\nabla_{\theta} E[f(\vec{x})]$ rather than fitting θ to elite

- Covariance^{Matrix} Adaptation

- similar to cross-entropy, but weight elite samples



Population Methods

Genetic Algorithms

Chromosome: vector of bits / real numbers

$[0, 1, 1, 0, \dots]$

$[5.2, 6.3, \dots]$ \vec{x}

\leftarrow in practice

Loop


Selection

Crossover

Mutation


Selection

Truncation




select best k individuals
randomly select from best k

Tournament



randomly select k
take best

Roulette Wheel




Crossover

Parent 1



Parent 2



Single Point
child




Two Point



Uniform



Mutation



Bitwise: each bit has an independent probability of flipping
Gaussian: each element has Gaussian noise added to it

Particle Swarm

initialize population $\{\vec{x}_i\}_{i=1}^m$

hyperparams: w, c_1, c_2

loop

for $i \in 1..m$

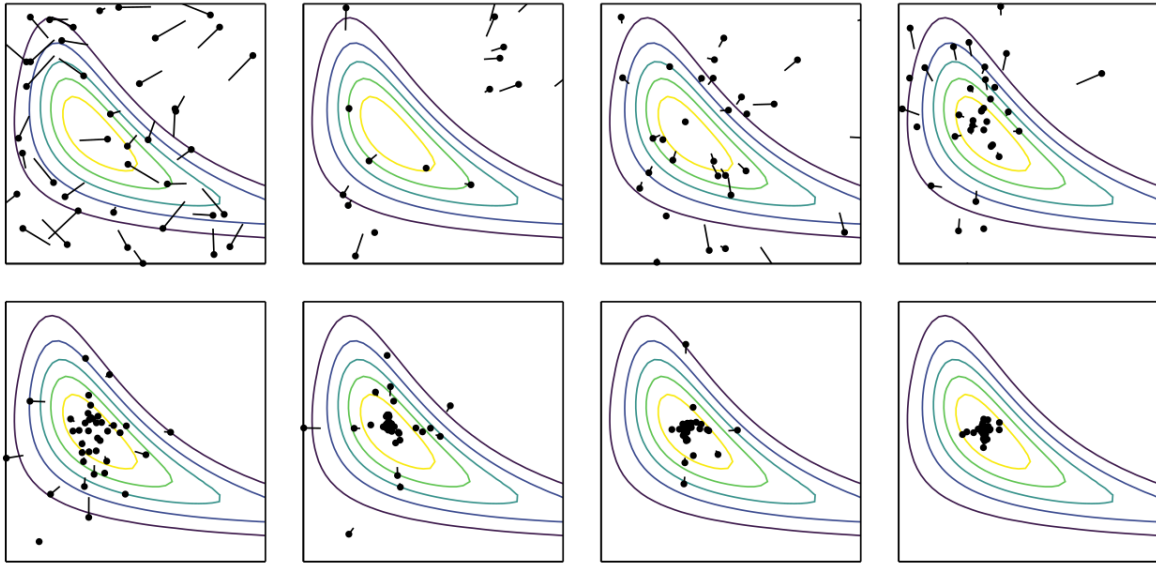
$$\vec{x}_i \leftarrow \vec{x}_i + \vec{v}_i$$

$$\vec{v}_i \leftarrow w \vec{v}_i + c_1 r_1 (\vec{x}_{best} - \vec{x}_i) + c_2 r_2 (\vec{x}_{best} - \vec{x}_i)$$

if $f(\vec{x}_i) < f(\vec{x}_{best})$
 $\vec{x}_{best} \leftarrow \vec{x}_i$

$\text{rand}(U(1,0))$

update \vec{x}_{best}



Firefly Algorithm

if $f(\vec{b}) < f(\vec{a})$
 \vec{a} moves toward \vec{b}

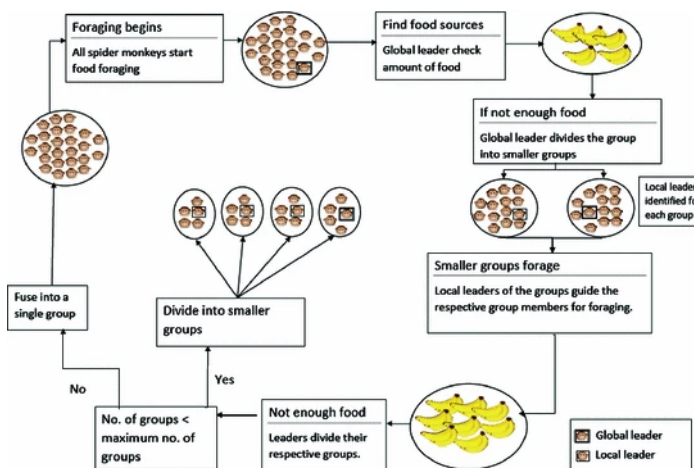

$$\vec{a} \leftarrow \vec{a} + \beta I(\|\vec{b} - \vec{a}\|) (\vec{b} - \vec{a}) + \alpha \epsilon$$

\swarrow intensity \nwarrow Gaussian
 \nwarrow hyperparameters

$$I(r) = \frac{1}{r^2}$$

$$I(r) = e^{-\gamma r}$$

$$I(r) = e^{-\gamma r^2}$$



§ These algorithms include the following: ant colony optimization, artificial bee colony algorithm, artificial fish swarm, artificial flora optimization algorithm, bacterial foraging optimization, bat algorithm, big bang-big crunch algorithm, biogeography-based optimization, bird mating optimizer, cat swarm optimization, cockroach swarm optimization, cuckoo search, design by shopping paradigm, dolphin echolocation algorithm, elephant herding optimization, firefly algorithm, flower pollination algorithm, fruit fly optimization algorithm, galactic swarm optimization, gray wolf optimizer, grenade explosion method, harmony search algorithm, hummingbird optimization algorithm, hybrid glowworm swarm optimization algorithm, imperialist competitive algorithm, intelligent water drops, invasive weed optimization, mine bomb algorithm, monarch butterfly optimization, moth-flame optimization algorithm, penguin search optimization algorithm, quantum-behaved particle swarm optimization, salp swarm algorithm, teaching-learning-based optimization, and whale optimization algorithm.

Why to use a meta heuristic alg. (nature inspired)

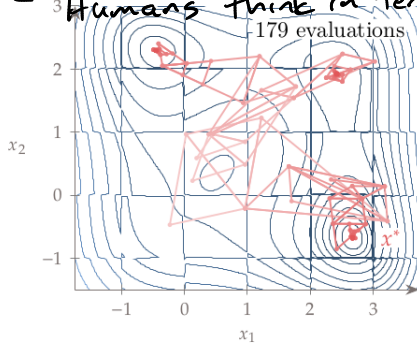
Pro

- There are NP-hard problems, usually involve many local minima
- Good for cheap function evaluation
- No need to be continuous/diff (useful for black-box)
- Highly Parallelizable
- Humans think in terms of stories

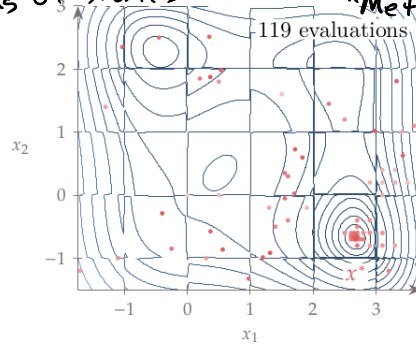
Con

- Few Guarantees
- Hard to know if a particular alg will work (hyperparameters)
- High population (bad if function eval is expensive)
- Poor scaling to high dimension (see below)
- Each new alg has own set of metaphors


"Meta-heuristics: the metaphor exposed"




Nelder-Mead algorithm




Generalized pattern search




DIRECT algorithm



Genetic algorithm



Particle swarm optimization



Quasi-Newton method




Fig. 7.1 Cost of optimization for increasing number of design variables in the n -dimensional Rosenbrock function. A gradient-free algorithm is compared with a gradient-based algorithm, with gradients computed analytically. The gradient-based algorithm has much better scalability.

Hybrid- Approach

stochastic / population \rightarrow explore new neighborhoods


descent \rightarrow find local minimum within neighborhood

$$\hat{f}(x) = f(x^{*2})$$




$$\hat{f}(x) = f(x^{*1})$$

Lamarckian



Baldwinian



Advantage: keeps points spread out